



Statistical Machine Learning

Lecture 5 – Cross-validation and the bias-variance trade-off



UPPSALA
UNIVERSITET

Andreas Lindholm

Division of Systems and Control
Department of Information Technology
Uppsala University

andreas.lindholm@it.uu.se

Summary of Lecture 4 (I/III)

Linear Discriminant Analysis (LDA) uses Bayes' theorem for classification, by writing the conditional class probabilities as,

$$\Pr(y = k | \mathbf{x}) = \frac{p(\mathbf{x} | y = k) \Pr(y = k)}{\sum_{j=1}^K p(\mathbf{x} | y = j) \Pr(y = j)}.$$

where

- $\pi_k = \Pr(y = k)$ is the *prior* probability of class k .
- $p(\mathbf{x} | y = k) = \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma})$ is the (assumed) probability density of \mathbf{x} for an observation that comes from the k th class.

The **parameters** are: $\pi_1, \dots, \pi_K, \boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_K, \boldsymbol{\Sigma}$

Summary of Lecture 4 (II/III)

The parameters are estimated as the class frequencies and (within class) sample means and covariances, respectively,

$$\begin{aligned}\hat{\pi}_k &= n_k/n & k &= 1, \dots, K, \\ \hat{\boldsymbol{\mu}}_k &= \frac{1}{n_k} \sum_{i:y_i=k} \mathbf{x}_i, & k &= 1, \dots, K, \\ \hat{\boldsymbol{\Sigma}} &= \frac{1}{n-K} \sum_{k=1}^K \sum_{i:y_i=k} (\mathbf{x}_i - \hat{\boldsymbol{\mu}}_k)(\mathbf{x}_i - \hat{\boldsymbol{\mu}}_k)^\top.\end{aligned}$$

The LDA classifier assigns a test input \mathbf{x}_\star to class k for which the predicted probability $\Pr(y = k \mid \mathbf{x}_\star)$, or equivalently,

$$\hat{\delta}_k(\mathbf{x}_\star) = \mathbf{x}_\star^\top \hat{\boldsymbol{\Sigma}}^{-1} \hat{\boldsymbol{\mu}}_k - \frac{1}{2} \hat{\boldsymbol{\mu}}_k^\top \hat{\boldsymbol{\Sigma}}^{-1} \hat{\boldsymbol{\mu}}_k + \log \hat{\pi}_k$$

is largest. Thus, LDA is a **linear classifier**.

Summary of Lecture 4 (III/III)

Non-parametric models are not specified using a fixed-dimensional vector of parameters. Instead, they allow the flexibility of the model to grow with the amount of available data.

One example is the **k -nearest neighbour classifier**.

Given training data $\mathcal{T} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, for a test input \mathbf{x}_\star ,


1. Identify the k training inputs \mathbf{x}_i nearest to \mathbf{x}_\star
2. Classify \mathbf{x}_\star according to a majority vote amongst these k training samples.



Evaluating a supervised machine learning method

Example

Let's say we have to design a subsystem for a self-driving car, which distinguishes between cyclists and pedestrians in camera images



- Input x : The color (RGB value) of each pixel
 - Output $y \in \{\text{cyclist}, \text{pedestrian}\}$
- a classification problem

Example

1. We spend a lot of resources to collect training data:
10 000 images of cyclists and pedestrians (each with a class label).
2. We train a classifier¹ using all training data. It classifies 99.7% of the training data correctly.
3. The system is integrated into the self-driving car, and eventually sold to a customer.

**How many cyclists and pedestrians will be classified correctly
when the car drives around in the real world?**

¹Perhaps a convolutional neural network, Lecture 9

Let's generalize using math

- Training data $\mathcal{T} \triangleq \{\mathbf{x}_i, y_i\}_{i=1}^n$
- Prediction $\hat{y}(\mathbf{x}; \mathcal{T})$, output from a method learned using \mathcal{T} ($\hat{g}(\mathbf{x})$ in previous lecture)
- Error function $E(\hat{y}, y) = \begin{cases} \mathbb{I}(y \neq \hat{y}) & \text{misclassification (classification problems)} \\ (y - \hat{y})^2 & \text{mean square error (regression problems)} \end{cases}$
- Training data error $E_{\text{train}} \triangleq \frac{1}{n} \sum_{i=1}^n E(\hat{y}(\mathbf{x}_i; \mathcal{T}), y_i)$
- The distribution of inputs and outputs in the real world $p(\mathbf{x}_\star, y_\star)$
- New data error $E_{\text{new}} \triangleq \int E(\hat{y}(\mathbf{x}_\star; \mathcal{T}), y_\star) p(\mathbf{x}_\star, y_\star) d\mathbf{x}_\star dy_\star \triangleq \mathbb{E}_\star [E(\hat{y}(\mathbf{x}_\star; \mathcal{T}), y_\star)]$.

(\mathbb{E}_\star = expected value over all future test inputs $\mathbf{x}_\star, y_\star$)

Connecting to the example


- \mathcal{T} : The 10 000 images
- $\hat{y}(\mathbf{x}; \mathcal{T})$: the prediction (cyclist/pedestrian) from the trained classifier
- $E(\hat{y}, y)$: whether the prediction for an image is correct or not
- $E_{\text{train}} = 0.003$ (how many cyclists and pedestrians it predicted wrong in the training data)
- $p(\mathbf{x}_\star, y_\star)$: The distribution of cyclists and pedestrians in the real world (appearance on camera and how often they occur)
- $E_{\text{new}} = ???$ (how many cyclists and pedestrians it will predict wrong in the real world)

What to be said about E_{new} ?

- A useful machine learning method gives a small E_{new} for our problem
- It would be useful to know E_{new} , so that we can tell if the method is good for our problem
- E_{new} depends on $p(\mathbf{x}, y)$, which we do not know

Let's try to *estimate* E_{new} !

Estimating E_{new} : test data approach



Set aside a part of the data and use only as **test data** for estimating E_{new} :

1. Use the training data \mathcal{T} for learning the model
2. Compute the predictions $\hat{y}(\mathbf{x}_*; \mathcal{T})$ for all data points \mathbf{x}_* in the test data set
3. Take the average of $E(\hat{y}(\mathbf{x}_*; \mathcal{T}), y_*)$ as an estimate of E_{new}

+ Simple

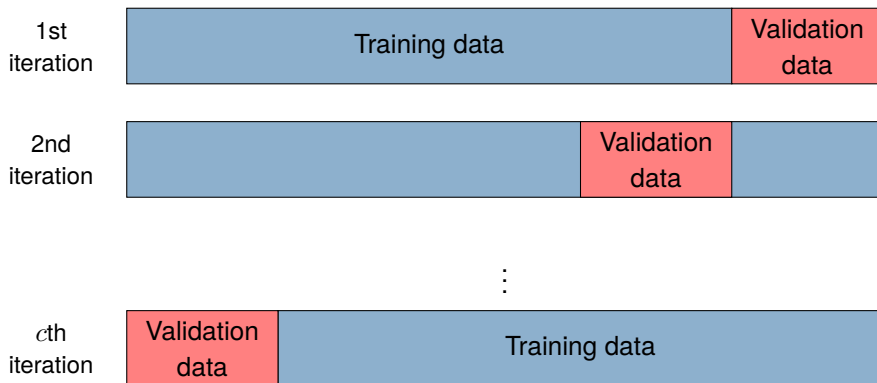
! To have a good estimate of E_{new} , the test dataset has to be large

! To get a good training of the model, the training dataset has to be large

– Not all data is used for learning

Always split **randomly** between training and test data!

Estimating E_{new} : c -fold cross-validation



Iterate the test data approach c times, and take the average as an estimate of E_{new} .

- + Gives a better estimate of E_{new}
- Computationally more demanding than the test data approach

Other flavors of cross-validation


- Leave-one-out cross-validation: c -fold cross-validation with $c = n$
- Monte Carlo cross-validation: random selection of validation set at each iteration

Example: cross-validation for model selection

All models are wrong, but some are useful. — George Box

Binary classification, $p = 2$. Should we use logistic regression, QDA or k -NN for this problem?


We would like to pick the method which has lowest E_{new} .
Since we only have the data, we can only *estimate* E_{new} .



	logistic regression	QDA	k -NN ($k = 1$)
Estimated* E_{new}	0.145	0.115	0.13

*using cross-validation

Example: cross-validation for model selection



	logistic regression	QDA	k -NN ($k = 1$)
Estimated E_{new}	0.15	0.12	0.13
E_{train}	0.14	0.11	0.0

Different methods have different E_{new} for different problems.

Why?


Model complexity

The more prone a model is to adapt to complicated pattern in the data, the higher the model complexity (or model flexibility):



It is possible to make a formal definition of model complexity. We will not do that in this course.

The example again



Generalization error


In general, we can say

$$E_{\text{new}} = E_{\text{train}} + \text{generalization error}$$


Typically,

- Higher model complexity \implies larger generalization error
- Higher model complexity \implies smaller E_{train}


Model complexity, E_{train} and E_{new}



Generalization error and training data size



Low model complexity



High model complexity

Digging deeper to understand E_{new}

Collecting training data is a random process (sampling \mathcal{T} from $p(\mathbf{x}, y)$) meaning there is randomness in the learned model.

Idea: Use probability theory to describe E_{new} in terms of **random effects in the model** (*variance*) and **systematic effects in the model** (*bias*).

Let's focus on the regression setting.

Bias-variance decomposition I/III

First, introduce \bar{E}_{new} where we average E_{new} over different training datasets,

$$\bar{E}_{\text{new}} \triangleq \mathbb{E}_{\mathcal{T}} [E_{\text{new}}]$$

where $\mathbb{E}_{\mathcal{T}}$ = expected value over all possible training data sets \mathcal{T} drawn from $p(\mathbf{x}, y)$.

Technically, it is \bar{E}_{new} (and not E_{new}) which we are going to decompose into bias and variance.

Now assume that "the real world" works as $y = f(\mathbf{x}) + \epsilon$, where ϵ is uncorrelated with \mathbf{x} and has mean 0

$$\begin{aligned} E_{\text{new}} &= \mathbb{E}_{\star} [E(\hat{y}(\mathbf{x}_{\star}; \mathcal{T}), y_{\star})] = \mathbb{E}_{\star} [(\hat{y}(\mathbf{x}_{\star}; \mathcal{T}) - f(\mathbf{x}_{\star}) - \epsilon)^2] \\ &= \mathbb{E}_{\star} [(\hat{y}(\mathbf{x}_{\star}; \mathcal{T}))^2] - 2\mathbb{E}_{\star} [\hat{y}(\mathbf{x}_{\star}; \mathcal{T})f(\mathbf{x}_{\star})] + 2\mathbb{E}_{\star} [f(\mathbf{x}_{\star}) - \hat{y}(\mathbf{x}_{\star}; \mathcal{T})] \underbrace{\mathbb{E}_{\star} [\epsilon]}_0 + \mathbb{E}_{\star} [(f(\mathbf{x}_{\star}))^2] + \underbrace{\mathbb{E}_{\star} [\epsilon^2]}_{\sigma^2} \\ &= \mathbb{E}_{\star} [(\hat{y}(\mathbf{x}_{\star}; \mathcal{T}))^2] - 2\hat{y}(\mathbf{x}_{\star}; \mathcal{T})f(\mathbf{x}_{\star}) + (f(\mathbf{x}_{\star}))^2 + \sigma^2 \end{aligned}$$

Bias-variance decomposition II/III

Introduce

$$g(\mathbf{x}_*) = \mathbb{E}_{\mathcal{T}} [\hat{y}(\mathbf{x}_*; \mathcal{T})],$$

the predicted output from the “average model”.

$$\begin{aligned}
 \bar{E}_{\text{new}} &= \mathbb{E}_{\mathcal{T}} [E_{\text{new}}] = \mathbb{E}_{\mathcal{T}} [\mathbb{E}_{\star} [(\hat{y}(\mathbf{x}_*; \mathcal{T}))^2 - 2\hat{y}(\mathbf{x}_*; \mathcal{T})f(\mathbf{x}_*) + (f(\mathbf{x}_*))^2] + \sigma^2] \\
 &= \mathbb{E}_{\star} \left[\mathbb{E}_{\mathcal{T}} [(\hat{y}(\mathbf{x}_*; \mathcal{T}))^2] - 2 \underbrace{\mathbb{E}_{\mathcal{T}} [\hat{y}(\mathbf{x}_*; \mathcal{T})]}_{g(\mathbf{x}_*)} f(\mathbf{x}_*) + (f(\mathbf{x}_*))^2 \right] + \sigma^2 \\
 &= \mathbb{E}_{\star} [\mathbb{E}_{\mathcal{T}} [(\hat{y}(\mathbf{x}_*; \mathcal{T}))^2] - (g(\mathbf{x}_*))^2 + (g(\mathbf{x}_*))^2 - 2g(\mathbf{x}_*)f(\mathbf{x}_*) + (f(\mathbf{x}_*))^2] + \sigma^2 \\
 &= \underbrace{\mathbb{E}_{\star} [\mathbb{E}_{\mathcal{T}} [(\hat{y}(\mathbf{x}_*; \mathcal{T}))^2] - (g(\mathbf{x}_*))^2]}_{\text{Variance}} + \underbrace{\mathbb{E}_{\star} [(g(\mathbf{x}_*) - f(\mathbf{x}_*))^2]}_{\text{Bias}^2} + \underbrace{\sigma^2}_{\text{Irreducible error}}
 \end{aligned}$$


Bias-variance decomposition III/III

Technical interpretation:

- **Bias²** $\mathbb{E}_\star \left[(g(\mathbf{x}_\star) - f(\mathbf{x}_\star))^2 \right]$: The part of \bar{E}_{new} that is due to the fact that, no matter how much training data is used, the model cannot represent the true f
- **Variance** $\mathbb{E}_\star \left[\mathbb{E}_{\mathcal{T}} \left[(\hat{y}(\mathbf{x}_\star; \mathcal{T}))^2 \right] - (g(\mathbf{x}_\star))^2 \right]$: The part of \bar{E}_{new} that is due to the variance in the training dataset

Intuitive interpretation:

- **Bias**: The inability of a method to describe the complicated patterns we would like it to describe. Low model complexity.
- **Variance**: How sensitive a method is to the training data. High model complexity.



Finding a balanced fit (neither over- nor underfit) is called the **the bias-variance tradeoff**.

Regression example


The data ($n = 10$) comes from

$$y = 5 - 2x + x^3 + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 1),$$


and our regression model is

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \varepsilon.$$



We use ridge regression to tune model complexity/bias-variance.



Regression example



Bias, variance and training data size



Make the tradeoff

Some methods for **decreasing the model complexity/increasing the bias/decreasing the variance**:

- Increase k in k -NN
- Regularization
- Bagging
- Early stopping (for methods trained using optimization, notably deep learning)
- Dropout (deep learning)

Warning! β_0 in linear regression (and later deep learning) is sometimes called “bias term”. That is **completely unrelated** to bias in this context.

Summary

- E_{new} = how good a method will perform “in production”
- We can *estimate* E_{new} well using cross-validation \rightarrow practical tool for model selection (choosing between models, k in k -NN, regularization parameter, ...)
- We have made two decompositions of E_{new} (or technically its average \bar{E}_{new}):
 - $\bar{E}_{\text{new}} = \bar{E}_{\text{train}} + \text{generalization error}$
 - $\bar{E}_{\text{new}} = \text{Bias}^2 + \text{Variance} + \text{Irreducible error}$

A few concepts to summarize lecture 5

$E(y, \hat{y})$: Error function which compares predictions \hat{y} to true output y : MSE for regression, misclassification for classification.

E_{train} : The training data error (E_{train} small = the method fits the training data well).

E_{new} : The expected new data error; how well a method will perform when faced with an endless stream of new data.

Cross-validation: A method for estimating E_{new} using the training data.

Model complexity: How prone a method is to adapt to complicated patterns in the training data.

Overfitting: When a given method yields a smaller E_{train} and larger E_{new} than a model with lower model complexity would have done. That happens because the method/model is capturing patterns in the training data caused by random chance rather than true properties of the underlying function.

Bias: The inability of a method to describe the true patterns in the classification or regression problem. Low model complexity.

Variance: Sensitivity to random effects (noise) in the training data. High model complexity.